

Pseudocode

Algorithm: NaiveBayesClassifier

Input:

X (training data)
y (class labels)

Output:

Model parameters (class priors, feature likelihoods)

Training Phase:

For each class c in y:
 $P(c) = \text{count}(c) / \text{total samples}$
 For each feature f:
 Compute $P(f | c)$

Prediction Phase:

For each test sample x:
 For each class c:
 $\text{score}(c) = P(c)$
 For each feature f in x:
 $\text{score}(c) = \text{score}(c) * P(f | c)$
 Predict class with maximum score

Python Code

```
import numpy as np

class NaiveBayes:
    def fit(self, X, y):
        self.classes = np.unique(y)
        self.mean = {}
        self.var = {}
        self.prior = {}

        for c in self.classes:
            X_c = X[y == c]
            self.mean[c] = np.mean(X_c, axis=0)
            self.var[c] = np.var(X_c, axis=0)
            self.prior[c] = X_c.shape[0] / X.shape[0]
```

```
def gaussian_prob(self, x, mean, var):
    eps = 1e-6
    coeff = 1 / np.sqrt(2 * np.pi * var + eps)
    exponent = np.exp(-(x - mean) ** 2 / (2 * var + eps))
    return coeff * exponent

def predict(self, X):
    predictions = []

    for x in X:
        posteriors = []
        for c in self.classes:
            prior = np.log(self.prior[c])
            likelihood = np.sum(
                np.log(self.gaussian_prob(x, self.mean[c], self.var[c])))
            posteriors.append(prior + likelihood)
        predictions.append(self.classes[np.argmax(posteriors)])

    return np.array(predictions)
```